



AtScale + Snowflake

Bridging Business Intelligence & Data Science in the Data Cloud

Today's Speakers





Simon FieldSnowcat Technical Director Snowflake

Simon works in Snowflakes
Customer Acceleration Team
(SnowCAT), supporting customers
to utilise new and advanced
product capabilities within
Snowflakes Data Cloud to improve
the value they derive from their
data.

Simon has worked in the field of Advanced Analytics, Data Warehousing, Big Data and Data Science for over 30 years, helping organisations make the transition to data-driven decision making.



Daniel GrayVP, Solutions Engineering AtScale

Daniel brings rich experience in technical solutions engineering as well as software engineering to his work with global enterprise organizations.

Prior to joining AtScale to lead the Solutions Engineering team, Daniel spent many years in the analytics space including Hewlett-Packard's Advanced Technology Center, Vertica, and Domino Data Lab.

SNOWFLAKE PLATFORM



DATA CONSUMERS

DATA MONETIZATION

OPERATIONAL REPORTING

AD HOC ANALYSIS

REAL-TIME ANALYTICS

Google Cloud





DATA

SOURCES

OLTP DATABASES

ENTERPRISE

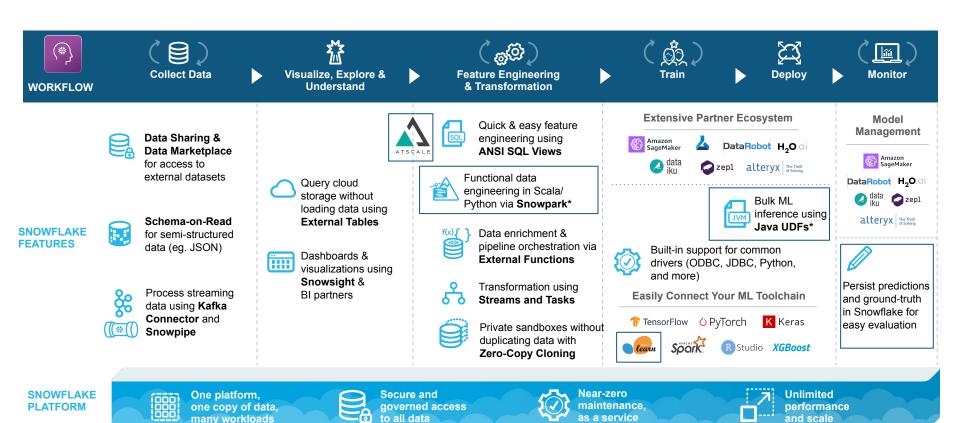
APPLICATIONS

THIRD-PARTY

WEB/LOG DATA

IoT

DATA SCIENCE WITH SNOWFLAKE

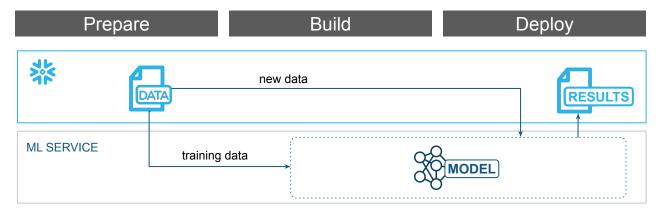


×

Model Inference : External Functions or Java UDFs

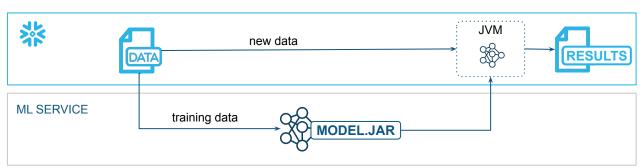
EXTERNAL SERVICE

Data continuously travels to externally hosted model (REST API) E.g. AWS Lambda



WITH JAVA UDF

Model packaged as java file (.jar) runs where data lives





ML partners with .JAR models: DataRobot, Dataiku, H2O or bring your own



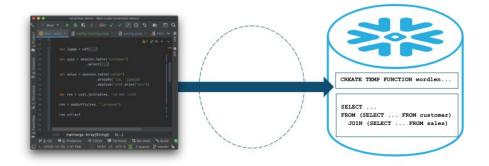
A new developer experience that allows you to write functional code and execute it directly within Snowflake

Example Use Cases:

- Data transformation
- Data preparation and feature engineering
- ML Scoring / Inference to operationalize ML models in data pipelines
- ELT systems
- Data apps

Allows coders to:

- Write in their preferred language and tools
- Easily complete and debug data pipelines with familiar constructs such as DataFrames, functions and use third-party libraries.
- Pushes all processing into Snowflake and eliminates the need to have other processing systems



Snowpark pushes all of its operations directly to Snowflake without the need for Spark or any other intermediary.

SNOWFLAKE JAVA FUNCTIONS

Transform and augment your data using custom logic running right next to your data, with no need to manage a separate service.

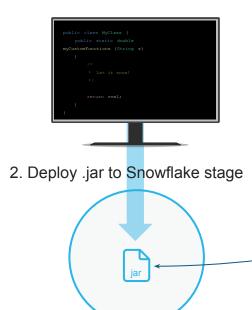
Example Scenarios:

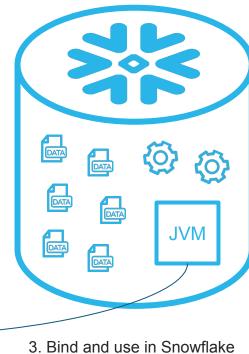
- ML Scoring
- Apply custom code
- Use third-party libraries

Benefits:

- Developers can build custom functionality in Snowflake using the JVM languages and popular libraries
- Snowpark 'publishes' functions developed in Scala as UDFs for execution in Snowflake via SQL or the Snowpark API.
- Users can access this functionality as if it were built in functions in Snowflake
- Administrators can rest easy: data never leaves Snowflake and access and execution permissions for functions can be controlled

1. Build with your tools

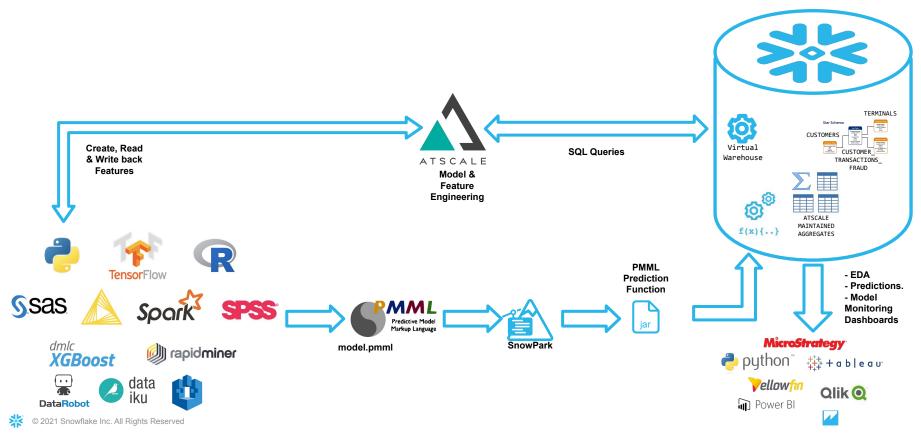




SNOWPARK + UDFs

```
Snowpark (Scala) Client or Scala Stored Procedure
val hasPII = udf(<PII detection code>)
  df = session.table("accident raw")
         .filter(hasPII("summary"))
         .select("summary")
  df.show()
                                                    SELECT summary
                                          JAR
                                                    FROM ( SELECT *
                                                           FROM ( SELECT * FROM (ACCIDENT RAW)
          CREATE TEMP FUNCTION hasPII...
                                                                 WHERE haspii("summary")
snowflake*
```

Model-Portability standards enable model inference & MLOps in Snowflake



DATA SCIENCE WITH SNOWFLAKE BEST PRACTICES



Enrich datasets using **Data**Marketplace for improved model accuracy



Use **Streams & Tasks** to build end-to-end ML pipelines



Create datasets without loading data into Snowflake via External Tables



Leverage External & JAVA Functions for training or to get predictions



Use **Zero-Copy Clones** for training snapshots



Views to create repository of ML features used for training and prediction



Optimize training instance memory usage by using **Snowflake SQL** for aggregation & sampling



Use **SnowPark** for functional programming with **dataframes** running in Snowflake

SUMMARY

- ☐ AtScale enables data, features and relationships to be modelled over Snowflake tables.
- Native Data Frame support via Snowpark enables Data Engineers and Scientists to build data engineering pipelines and execute models.
- ☐ Model storage/persistence and interoperability via PMML (and other) open model format.
- Java UDF allows fast compiled custom code execution within Snowflake.
- Access to Java based languages and libraries directly in Snowflake.