
















































































































 1st generation	 2nd generation	 2nd generation	 2nd generation	FIREBOLT 3rd generation
ARCHITECTURE					
Elasticity - separation of storage and compute	 Redshift RA3, Spectrum only	 Yes (query engine, no storage)	 Yes	 Yes	 Yes
Supported cloud infrastructure	 AWS only	 AWS only	 Google Cloud only	 AWS, Azure, Google Cloud	 AWS only
Isolated tenancy - option for dedicated resources	 Isolated tenant, use own VPC	 Multi-tenant pooled resources	 Multi-tenant on demand and reserved resources only	 Multi-tenant resources (Only VPS dedicated on isolated Snowflake account)	 Multi-tenant dedicated resources
Compute - node types	 Any size, and types	 N/A no choice	 No choice over (fixed) slot size	 1-128 nodes, unknown types	 1-128 choice of any types
Data - internal/external, writable storage	 Internal writable, external (Spectrum)	 External only	 External tables supported - 4 concurrent queries by default*	 External tables supported	 External tables (used for ingestion)
Security - data	 Isolated/VPC tenant for storage and compute Encryption at rest, RBAC	 Shared resources	 Separate customer keys, column-level encryption, encryption at rest, AEAD individual value encryption, RBAC	 Separate customer keys (only VPS is isolated tenant) Encryption at rest, RBAC	 Dedicated resources for storage and compute Encryption at rest, RBAC
Security - network	 Firewall, SSL, PrivateLink whitelist/blacklist control, isolated/VPC tenancy	 Shared resources	 TLS, Firewall (Google Cloud), TLS, VPN, whitelist/blacklist control part of GCP	 Firewall, SSL, PrivateLink whitelist/blacklist control, isolated for VPS only	 Firewall and WAF, SSL, PrivateLink whitelist/blacklist control

SCALABILITY					
Elasticity - (individual) query scalability	● Manual	● Automatic (but shared resources)	● Automatic allocation of each query to on demand, or reserved and flex slots	● 1-click cluster resize, no choice of node size	● 1-click cluster resize of node type, number of nodes
Elasticity - user (concurrent query) scalability	● Autoscale up to 10 clusters, 15 queries per cluster, 50 queued queries total, max.	● Limited - 20 concurrent queries by default	● Limited to 100 concurrent users by default*	● Autoscale up to 10 warehouses. Limited to 20 DML writes in queue per table	● Unlimited manual scaling
Write scalability - batch	● 1 master cluster	● N/A	● 1,500 load jobs/day (~1 per minute), 100,000 per project, 15TB per job, 6 hour max time*	● Strong	● Strong. Multi-master parallel batch
Write scalability - continuous	● Limited (table-level locking)	● N/A (mostly batch-centric storage)	● 1GB/sec w/ no dedup, 100MB./sec w/ dedup, 100K rows per second per table, 100K-500K per project by default*	● Limited to 20 DML writes in queue per table. 1 minute or greater ingestion latency recommended	● Multi-master continuous writes
Data scalability	● Only with RA3: 128 RA3 nodes, 8PB of data.	● Up to 100 partitions per table, 100 buckets (default)	● No real limit	● No specified storage limit. 4XL data warehouse (128 nodes)	● No limit

PERFORMANCE					
Indexes	 None	 None	 None	 None	 Indices for data access, joins, aggregation, search
Query optimization - performance	 Limited cost-based optimization	 Limited cost-based optimization	 Cost-based optimization	 Cost-based optimization, vectorization	 Index- and cost-based optimization, vectorization JIT, pushdown optimization
Tuning	 Choice of (limited) node types	 No choice of resources	 Can only purchase reserved or flex slots	 Can only choose warehouse size not node types	 Choice of node types, Indexing
Storage format	 Native Redshift storage (not Spectrum)	 S3	 Optimized (capacitor) on Colossus	 Optimized micro-partition storage (S3), separate RAM	 Optimized F3 storage (on S3). Data access integrated across disk, SSD, RAM)
Ingestion performance	 Batch-centric (table-level locking)	 N/A (storage and ingestion separate)	 Writes 1 row at a time. Limits of 100K messages/sec by default*	 Batch-centric (micro-partition level locking, limit of 20 queued writes per table)	 Multi-master, lock-free high performance ingestion with unlimited scale for batch or continuous ingestion
Ingestion latency	 Batch-centric (minute-level)	 Not well suited for low latency visibility since unable to see as new values during ingestion	 Immediately visible during ingestion	 Batch write preferred (1+ minute interval). Requires rewrite of entire micro-partition	 Immediately visible during ingestion
Partitioning	 Distribution, sort keys	 Partition pruning	 Partitions, pruning	 Micro-partition / pruning, cluster keys	 F3 with sparse indexing
Caching	 Result cache	 None	 Result cache (24 hours), shared memory	 Result cache, materialized view	 F3 (cache) with aggregating, join, and search indexes
Semi-structured data - native JSON functions within SQL	 Limited	 Yes (Lambda)	 Yes	 Yes	 Yes (Lambda)
Semi-structured data - native JSON storage type	 No	 No	 Can store as strings or STRUCT. But requires UDFs for compute	 Limited VARIANT (single field)	 Yes (Nested array type, compact)
Semi-structured data - performance	 Slow (flattens JSON into table)	 Slow (full load into RAM, full scan)	 JSON strings slow. Can store as STRUCT and use UDF (JS)	 Slow (can require full load into RAM, full scan)	 Fast (array operations without full scans)

USE CASES					
Reporting	 Yes	 Yes	 Yes	 Yes	 Yes
Dashboards	 Fixed view	 Fixed view	 Fixed view	 Fixed view	 Fixed view, dynamic / changing data
Ad hoc, interactive analytics	 Seconds- minutes first- time query performance	 Seconds-minutes first-time query performance	 Seconds- minutes first- time query performance	 Seconds- minutes first- time query performance	 Sub-second to seconds first-time query performance
Operational or customer-facing analytics (high concurrency, continuously updating / streaming data)	 Slower query performance. Limited continuous writes and concurrency, limited semi- structured data support	 Slow query performanceand limited scale. Limited continuous writes and concurrency, slow semi- structured data performance	 Slower query performance. Limited to 100K continuous writes/table, 100 concurrent users by default*	 Slower query performance. Limited continuous writes and concurrency, slow semi- structured data performance	 Yes. Support continuous writes at scale, fast semi- structured performance
Data processing engine (Exports or publishes data)	 Unload data as Parquet	 Exports query results	 1GB max export file size, exports to Google cloud only*	 Export query results or table	 Export query results
Data science/ML	 Invoke ML (SageMaker) in SQL	 Export query results	 BigQuery ML	 Spark, Arrow, Python connectors, integration with ML tools, export query results	 Export query results

COST					
Administration - deployment, management	 Easy to provision, harder to configure and manage	 No administration or tuning	 No administration or tuning	 Easy to deploy and resize Strong performance visibility limited tuning	 Easy to deploy and resize Easy to add indexing, change node types
Choice - provision different cluster types on same data	 No	 No	 Yes	 Choice of fixed size warehouses	 Choice of node types, engine sizes
Choice - provision different number of nodes	 Yes	 No	 Up to 2000 flex (on demand) slots, Purchase reserved or flex slots 100 at a time with no limits	 Yes	 Yes
Choice - provision different node types	 Yes (Limited)	 No	 No	 No	 Yes
Pricing - compute	 \$0.25-13 per node on demand <40% less when up front	 None	 On demand - \$5/TB data processed, Flex slots \$4 for 100 slots per hour \$1700/month per 100 slots	 \$2-\$4+ per node. Fast analytics need large cluster (\$16-32+/hour) or greater	 Choose any any node. Compute costs range \$1-10/hour with 10x price-performance advantage
Pricing - storage	 Stored data only. RA3: \$24 per TB per month S3: AWS S3 costs	 N/A (not part of Athena)	 \$20/TB active storage, \$10/TB inactive	 All storage. \$23/40 per TB per month on demand/up front	 All storage. \$23/40 per TB per month on demand/up front
Pricing - transferred data	 Spectrum: \$5 per TB scanned (10MB min per query)	 \$5 per TB scanned (10MB per query)	 Batch is free. Streaming ingest \$0.01 per 200MB (\$50/TB), streaming reads \$1.1/TB	 None	 None