

Social Media Mining

An Introduction

REZA ZAFARANI

MOHAMMAD ALI ABBASI

HUAN LIU



Draft Version: April 20, 2014

By permission of Cambridge University Press, this preprint is free. Users can make one hardcopy for personal use, but not for further copying or distribution (either print or electronically). Users may link freely to the book's website: <http://dmml.asu.edu/smm>, but may not post this preprint on other web sites.

SOCIAL MEDIA MINING

Contents

1	Introduction	15
1.1	What is Social Media Mining	16
1.2	New Challenges for Mining	16
1.3	Book Overview and Reader's Guide	18
1.4	Summary	21
1.5	Bibliographic Notes	22
1.6	Exercises	24
I	Essentials	27
2	Graph Essentials	29
2.1	Graph Basics	30
2.1.1	Nodes	30
2.1.2	Edges	31
2.1.3	Degree and Degree Distribution	32
2.2	Graph Representation	35
2.3	Types of Graphs	37
2.4	Connectivity in Graphs	39
2.5	Special Graphs	44
2.5.1	Trees and Forests	44
2.5.2	Special Subgraphs	45
2.5.3	Complete Graphs	47
2.5.4	Planar Graphs	47
2.5.5	Bipartite Graphs	48

2.5.6	Regular Graphs	48
2.5.7	Bridges	49
2.6	Graph Algorithms	49
2.6.1	Graph/Tree Traversal	49
2.6.2	Shortest Path Algorithms	53
2.6.3	Minimum Spanning Trees	56
2.6.4	Network Flow Algorithms	57
2.6.5	Maximum Bipartite Matching	61
2.6.6	Bridge Detection	64
2.7	Summary	66
2.8	Bibliographic Notes	68
2.9	Exercises	69
3	Network Measures	73
3.1	Centrality	74
3.1.1	Degree Centrality	74
3.1.2	Eigenvector Centrality	75
3.1.3	Katz Centrality	78
3.1.4	PageRank	80
3.1.5	Betweenness Centrality	82
3.1.6	Closeness Centrality	84
3.1.7	Group Centrality	85
3.2	Transitivity and Reciprocity	87
3.2.1	Transitivity	87
3.2.2	Reciprocity	91
3.3	Balance and Status	92
3.4	Similarity	95
3.4.1	Structural Equivalence	96
3.4.2	Regular Equivalence	98
3.5	Summary	101
3.6	Bibliographic Notes	102
3.7	Exercises	103
4	Network Models	105
4.1	Properties of Real-World Networks	106
4.1.1	Degree Distribution	106
4.1.2	Clustering Coefficient	109
4.1.3	Average Path Length	109

4.2	Random Graphs	110
4.2.1	Evolution of Random Graphs	112
4.2.2	Properties of Random Graphs	115
4.2.3	Modeling Real-World Networks with Random Graphs	118
4.3	Small-World Model	119
4.3.1	Properties of the Small-World Model	121
4.3.2	Modeling Real-World Networks with the Small-World Model	124
4.4	Preferential Attachment Model	125
4.4.1	Properties of the Preferential Attachment Model . . .	126
4.4.2	Modeling Real-World Networks with the Preferen- tial Attachment Model	128
4.5	Summary	129
4.6	Bibliographic Notes	131
4.7	Exercises	132
5	Data Mining Essentials	135
5.1	Data	137
5.1.1	Data Quality	141
5.2	Data Preprocessing	142
5.3	Data Mining Algorithms	144
5.4	Supervised Learning	144
5.4.1	Decision Tree Learning	145
5.4.2	Naive Bayes Classifier	148
5.4.3	Nearest Neighbor Classifier	150
5.4.4	Classification with Network Information	151
5.4.5	Regression	154
5.4.6	Supervised Learning Evaluation	157
5.5	Unsupervised Learning	159
5.5.1	Clustering Algorithms	160
5.5.2	Unsupervised Learning Evaluation	162
5.6	Summary	166
5.7	Bibliographic Notes	167
5.8	Exercises	169

II	Communities and Interactions	173
6	Community Analysis	175
6.1	Community Detection	179
6.1.1	Community Detection Algorithms	179
6.1.2	Member-Based Community Detection	181
6.1.3	Group-Based Community Detection	188
6.2	Community Evolution	197
6.2.1	How Networks Evolve	198
6.2.2	Community Detection in Evolving Networks	201
6.3	Community Evaluation	204
6.3.1	Evaluation with Ground Truth	204
6.3.2	Evaluation without Ground Truth	209
6.4	Summary	211
6.5	Bibliographic Notes	212
6.6	Exercises	214
7	Information Diffusion in Social Media	217
7.1	Herd Behavior	220
7.1.1	Bayesian Modeling of Herd Behavior	222
7.1.2	Intervention	224
7.2	Information Cascades	225
7.2.1	Independent Cascade Model (ICM)	226
7.2.2	Maximizing the Spread of Cascades	229
7.2.3	Intervention	231
7.3	Diffusion of Innovations	232
7.3.1	Innovation Characteristics	232
7.3.2	Diffusion of Innovations Models	233
7.3.3	Modeling Diffusion of Innovations	236
7.3.4	Intervention	239
7.4	Epidemics	240
7.4.1	Definitions	242
7.4.2	SI Model	243
7.4.3	SIR Model	245
7.4.4	SIS Model	246
7.4.5	SIRS Model	248
7.4.6	Intervention	249
7.5	Summary	251

7.6	Bibliographic Notes	252
7.7	Exercises	254

III Applications 257

8	Influence and Homophily	259
8.1	Measuring Assortativity	261
8.1.1	Measuring Assortativity for Nominal Attributes . . .	262
8.1.2	Measuring Assortativity for Ordinal Attributes . . .	265
8.2	Influence	268
8.2.1	Measuring Influence	268
8.2.2	Modeling Influence	273
8.3	Homophily	278
8.3.1	Measuring Homophily	278
8.3.2	Modeling Homophily	278
8.4	Distinguishing Influence and Homophily	280
8.4.1	Shuffle Test	280
8.4.2	Edge-Reversal Test	281
8.4.3	Randomization Test	282
8.5	Summary	285
8.6	Bibliographic Notes	286
8.7	Exercises	287
9	Recommendation in Social Media	289
9.1	Challenges	290
9.2	Classical Recommendation Algorithms	292
9.2.1	Content-Based Methods	292
9.2.2	Collaborative Filtering (CF)	293
9.2.3	Extending Individual Recommendation to Groups of Individuals	300
9.3	Recommendation Using Social Context	303
9.3.1	Using Social Context Alone	303
9.3.2	Extending Classical Methods with Social Context . .	304
9.3.3	Recommendation Constrained by Social Context . . .	306
9.4	Evaluating Recommendations	309
9.4.1	Evaluating Accuracy of Predictions	309
9.4.2	Evaluating Relevancy of Recommendations	310

9.4.3	Evaluating Ranking of Recommendations	311
9.5	Summary	314
9.6	Bibliographic Notes	315
9.7	Exercises	316
10	Behavior Analytics	319
10.1	Individual Behavior	319
10.1.1	Individual Behavior Analysis	321
10.1.2	Individual Behavior Modeling	326
10.1.3	Individual Behavior Prediction	327
10.2	Collective Behavior	332
10.2.1	Collective Behavior Analysis	332
10.2.2	Collective Behavior Modeling	337
10.2.3	Collective Behavior Prediction	338
10.3	Summary	340
10.4	Bibliographic Notes	341
10.5	Exercises	343
	Bibliography	345
	Index	373

Preface

We live in an age of big data. With hundreds of millions of people spending countless hours on social media to share, communicate, connect, interact, and create user-generated data at an unprecedented rate, social media has become one unique source of big data. This novel source of rich data provides unparalleled opportunities and great potential for research and development. Unfortunately, more data does not necessarily beget more good, only more of the right (or relevant) data that enables us to glean gems. Social media data differs from traditional data we are familiar with in data mining. Thus, new computational methods are needed to mine the data. Social media data is noisy, free-format, of varying length, and multimedia. Furthermore, social relations among the entities, or social networks, form an inseparable part of social media data; hence, it is important that social theories and research methods be employed with statistical and data mining methods. It is therefore a propitious time for social media mining.

Social media mining is a rapidly growing new field. It is an interdisciplinary field at the crossroad of disparate disciplines deeply rooted in computer science and social sciences. There are an active community and a large body of literature about social media. The fast-growing interests and intensifying need to harness social media data require research and the development of tools for finding insights from big social media data. This book is one of the intellectual efforts to answer the novel challenges of social media. It is designed to enable students, researchers, and practitioners to acquire fundamental concepts and algorithms for social media mining.

Researchers in this emerging field are expected to have knowledge in different areas, such as data mining, machine learning, text mining, social network analysis, and information retrieval, and are often required to consult research papers to learn the state of the art of social media mining. To mitigate such a strenuous effort and help researchers get up to speed in a convenient way, we take advantage of our teaching and research of many years to survey, summarize, filter, categorize, and connect disparate research findings and fundamental concepts of social media mining. This book is our diligent attempt to provide an easy reference or entry point to help researchers quickly acquire a wide range of essentials of social media mining. Social media not only produces big user-generated data; it also has a huge potential for social science research, business development, and understanding human and group behavior. If you want to share a piece of information or a site on social media, you would like to grab precious attention from other equally eager users of social media; if you are curious to know what is hidden or who is influential in the complex world of social media, you might wonder how one can find this information in big and messy social media; if you hope to serve your customers better in social media, you certainly want to employ effective means to understand them better. These are just some scenarios in which social media mining can help. If one of these scenarios fits your need or you simply wish to learn something interesting in this emerging field of social media mining, this book is for you. We hope this book can be of benefit to you in accomplishing your goals of dealing with big data of social media.

Book Website and Resources

The book's website and further resources can be found at

<http://dmml.asu.edu/smm>

The website provides lecture slides, homework and exam problems, and sample projects, as well as pointers to useful material and resources that are publicly available and relevant to social media mining.

To the Instructors

The book is designed for a one-semester course for senior undergraduate or graduate students. Though it is mainly written for students with a

background in computer science, readers with a basic understanding of probability, statistics, and linear algebra will find it easily accessible. Some chapters can be skipped or assigned as a homework assignment for reviewing purposes if students have knowledge of a chapter. For example, if students have taken a data mining or machine learning course, they can skip Chapter 5. When time is limited, Chapters 6–8 should be discussed in depth, and Chapters 9 and 10 can be either discussed briefly or assigned as part of reading material for course projects.

Reza Zafarani
Mohammad Ali Abbasi
Huan Liu
Tempe, AZ
August, 2013

Acknowledgments

In the past several years, enormous pioneering research has been performed by numerous researchers in the interdisciplinary fields of data mining, social computing, social network analysis, network science, computer science, and the social sciences. We are truly dwarfed by the depth, breadth, and extent of the literature, which not only made it possible for us to complete a text on this emerging topic – *social media mining* – but also made it a seemingly endless task. In the process, we have been fortunate in drawing inspiration and obtaining great support and help from many people to whom we are indebted.

We would like to express our tremendous gratitude to the current and former members of the Data Mining and Machine Learning laboratory at Arizona State University (ASU); in particular, Nitin Agrawal, Salem Alelyani, Geoffrey Barbier, William Cole, Zhuo Feng, Magdiel Galan-Oliveras, Huiji Gao, Pritam Gundecha, Xia (Ben) Hu, Isaac Jones, Shamanth Kumar, Fred Morstatter, Sai Thejasvee Moturu, Ashwin Rajadesingan, Suhas Ranganath, Jiliang Tang, Lei Tang, Xufei Wang, and Zheng Zhao. Without their impressive accomplishments and continuing strides in advancing research in data mining, machine learning, and social computing, this book would have not been possible. Their stimulating thoughts, creative ideas, friendly aggressiveness, willingness to extend the research frontier, and cool company during our struggling moments (Arizona could be scorchingly hot in some months), directly and indirectly, offered us encouragement, drive, passion, and ideas, as well as critiques in the process toward the completion of the book.

This book project stemmed from a course on social computing offered in 2008 at ASU. It was a seminar course that enjoyed active participation by graduate students and bright undergraduates with intelligent and provocative minds. Lively discussion and heated arguments were fixtures of the seminar course. Since then, it has become a regular course, evolving into a focused theme on *social media mining*. Teaching assistants, students, and guest speakers in these annual courses were of significant help to us in choosing topics to include, determining the depth and extent of each topic, and offering feedback on lecture materials such as homework problems, slides, course projects, and reading materials.

We would like to especially thank Denny Abraham Cheriyan, Nitin Ahuja, Amy Baldwin, Sai Prasanna Baskaran, Gaurav Pandey, Prerna

Satija, Nitesh Kedia, Bernard Ma, Dhanyatha Manjunath, Girish Kumar Reddy Marthala, Apurv Patki, Gregory Scherer, Nikhil Sunkesula Bhaskar, Yajin Wang, and Ruozhou Yu for their detailed comments on the drafts of this book. In addition, we owe our gratitude to Daria Bazzi and Michael Meeder for their help in proofreading earlier versions of the book, Farzad Zafarani for preparing the book's website, Subbarao Kambhampati for reading an earlier draft and offering encouragement, and Rebecca Goolsby for continually challenging us in understanding social media and developing social computing tools for real-world applications of humanitarian assistance and disaster relief.

The idea of having this book published by Cambridge University Press began with a casual conversation with K. Selcuk Candan, a well-received Cambridge author. It was an enjoyable and pleasant process working with Cambridge University Press. We would like to thank Lauren Cowles, a senior editor of Mathematics and Computer Sciences at Cambridge, for her patience and kind support during the process, and the Cambridge team, David Jou and Joshua Penney, as well as Adrian Pereira and his colleagues at Aptara for their work on the production of the book.

Our research on social computing, data mining, and social network analysis has been, in part, supported by the Office of Naval Research, National Science Foundation, and Army Research Office.

We have truly enjoyed our collaboration in this arduous journey. We will certainly miss our weekly meetings and many missed deadlines.

Chapter 1

Introduction

With the rise of *social media*, the web has become a vibrant and lively realm in which billions of individuals all around the globe interact, share, post, and conduct numerous daily activities. Information is collected, curated, and published by *citizen journalists* and simultaneously shared or consumed by thousands of individuals, who give spontaneous feedback. Social media enables us to be connected and interact with each other anywhere and anytime – allowing us to observe human behavior in an unprecedented scale with a new lens. This social media lens provides us with golden opportunities to understand individuals at scale and to mine human behavioral patterns otherwise impossible. As a byproduct, by understanding individuals better, we can design better computing systems tailored to individuals’ needs that will serve them and society better. This new social media world has no geographical boundaries and incessantly churns out oceans of data. As a result, we are facing an exacerbated problem of big data – “drowning in data, but thirsty for knowledge.” Can data mining come to the rescue?

Social Media

Citizen
Journalism

Unfortunately, social media data is significantly different from the traditional data that we are familiar with in data mining. Apart from enormous size, the mainly user-generated data is noisy and unstructured, with abundant social relations such as friendships and followers-followees. This new type of data mandates new computational data analysis approaches that can combine social theories with statistical and data mining methods. The pressing demand for new techniques ushers in and entails a new interdisciplinary field – social media mining.

1.1 What is Social Media Mining

Social Atom

Social Molecule

Social media shatters the boundaries between the real world and the virtual world. We can now integrate social theories with computational methods to study how individuals (also known as *social atoms*) interact and how communities (i.e., *social molecules*) form. The uniqueness of social media data calls for novel data mining techniques that can effectively handle user-generated content with rich social relations. The study and development of these new techniques are under the purview of social media mining, an emerging discipline under the umbrella of data mining. ***Social Media Mining*** is the process of representing, analyzing, and extracting actionable patterns from social media data.

Social Media Mining

Social Media Mining, introduces basic concepts and principal algorithms suitable for investigating massive social media data; it discusses theories and methodologies from different disciplines such as computer science, data mining, machine learning, social network analysis, network science, sociology, ethnography, statistics, optimization, and mathematics. It encompasses the tools to formally represent, measure, model, and mine meaningful patterns from large-scale social media data.

Data Scientist

Social media mining cultivates a new kind of *data scientist* who is well versed in social and computational theories, specialized to analyze recalcitrant social media data, and skilled to help bridge the gap from what we know (social and computational theories) to what we want to know about the vast social media world with computational tools.

1.2 New Challenges for Mining

Social media mining is an emerging field where there are more problems than ready solutions. Equipped with interdisciplinary concepts and theories, fundamental principles, and state-of-the-art algorithms, we can stand on the shoulders of the giants and embark on solving challenging problems and developing novel data mining techniques and scalable computational algorithms. In general, social media can be considered a world of social atoms (i.e., individuals), entities (e.g., content, sites, networks, etc.), and interactions between individuals and entities. Social theories and social norms govern the interactions between individuals and entities. For effective social media mining, we collect information about individuals and

entities, measure their interactions, and discover patterns to understand human behavior.

Mining social media data is the task of mining user-generated content with social relations. This data¹ presents novel challenges encountered in social media mining.

Big Data Paradox. Social media data is undoubtedly big. However, when we zoom into individuals for whom, for example, we would like to make relevant recommendations, we often have little data for each specific individual. We have to exploit the characteristics of social media and use its multidimensional, multisource, and multisite data to aggregate information with sufficient statistics for effective mining.

Big Data
Paradox

Obtaining Sufficient Samples. One of the commonly used methods to collect data is via application programming interfaces (APIs) from social media sites. Only a limited amount of data can be obtained daily. Without knowing the population's distribution, how can we know that our samples are reliable representatives of the full data? Consequently, how can we ensure that our findings obtained from social media mining are any indication of true patterns that can benefit our research or business development?

Obtaining
Sufficient
Samples

Noise Removal Fallacy. In classic data mining literature, a successful data mining exercise entails extensive data preprocessing and noise removal as "garbage in and garbage out." By its nature, social media data can contain a large portion of noisy data. For this data, we notice two important observations: (1) blindly removing noise can worsen the problem stated in the big data paradox because the removal can also eliminate valuable information, and (2) the definition of noise becomes complicated and relative because it is dependent on our task at hand.

Noise
Removal
Fallacy

Evaluation Dilemma. A standard procedure of evaluating patterns in data mining is to have some kind of ground truth. For example, a dataset can be divided into training and test sets. Only the training data is used in learning, and the test data serves as ground truth for testing. However, ground truth is often not available in social media mining. Evaluating patterns from social media mining poses a seemingly insurmountable challenge. On the other hand, without credible evaluation, how can we guarantee the

Evaluation
Dilemma

¹The data has a power-law distribution and more often than not, data is not independent and identically distributed (i.i.d.) as generally assumed in data mining.

validity of the patterns?

This book contains basic concepts and fundamental principles that will help readers contemplate and design solutions to address these challenges intrinsic to social media mining.

1.3 Book Overview and Reader's Guide

This book consists of three parts. Part I, *Essentials*, outlines ways to represent social media data and provides an understanding of fundamental elements of social media mining. Part II, *Communities and Interactions*, discusses how communities can be found in social media and how interactions occur and information propagates in social media. Part III, *Applications*, offers some novel illustrative applications of social media mining. Throughout the book, we use examples to explain how things work and to deepen the understanding of abstract concepts and profound algorithms. These examples show in a tangible way how theories are applied or ideas are materialized in discovering meaningful patterns in social media data.

Consider an online social networking site with millions of members in which members have the opportunity to befriend one another, send messages to each other, and post content on the site. Facebook, LinkedIn, and Twitter are exemplars of such sites. To make sense of data from these sites, we resort to social media mining to answer corresponding questions. In Part I: Essentials (Chapters 2–5), we learn to answer questions such as the following:

1. Who are the most important people in a social network?
2. How do people befriend others?
3. How can we find interesting patterns in user-generated content?

These essentials come into play in Part II: Communities and Interactions (Chapters 6 and 7) where we attempt to analyze how communities are formed, how they evolve, and how the qualities of detected communities are evaluated. We show ways in which information diffusion in social media can be studied. We aim to answer general questions such as the following:

1. How can we identify communities in a social network?

2. When someone posts an interesting article on a social network, how far can the article be transmitted in that network?

In Part III: Application (Chapters 8–10), we exemplify social media mining using real-world problems in dealing with social media: measuring influence, recommending in a social environment, and analyzing user behavior. We aim to answer these questions:

1. How can we measure the influence of individuals in a social network?
2. How can we recommend content or friends to individuals online?
3. How can we analyze the behavior of individuals online?

To provide an overall picture of the book content, we created a dependency graph among chapters (Fig. 1.1) in which arrows suggest dependencies between chapters. Based on the dependency graph, therefore, a reader can start with Chapter 2 (graph essentials), and it is recommended that he or she read Chapters 5 (data mining essentials) and 8 (influence and homophily) before Chapter 9 (recommendation in social media). We have also color-coded chapter boxes that are of the same level of importance and abstraction. The darkest chapters are the essentials of this book, and the lightest boxes are those chapters that are more applied and have materials that are built on the foundation of other chapters.

Who Should Read This Book?

A reader with a basic computer science background and knowledge of data structures, search, and graph algorithms will find this book easily accessible. Limited knowledge of linear algebra, calculus, probability, and statistics will help readers understand technical details with ease. Having a data mining or machine learning background is a plus, but not necessary.

The book is designed for senior undergraduate and graduate students. It is organized in such a way that it can be taught in one semester to students with a basic prior knowledge of statistics and linear algebra. It can also be used for a graduate seminar course by focusing on more advanced chapters with the supplement of detailed bibliographical notes. Moreover, the book can be used as a reference book for researchers, practitioners, and project managers of related fields who are interested in both learning the




Figure 1.1: Dependency between Book Chapters. Arrows show dependencies and colors represent book parts.

basics and tangible examples of this emerging field and understanding the potentials and opportunities that social media mining can offer.

1.4 Summary

As defined by Kaplan and Haenlein [141], social media is the “group of internet-based applications that build on the ideological and technological foundations of Web 2.0, and that allow the creation and exchange of user-generated content.” There are many categories of social media including, but not limited to, social networking (Facebook or LinkedIn), microblogging (Twitter), photo sharing (Flickr, Photobucket, or Picasa), news aggregation (Google reader, StumbleUpon, or Feedburner), video sharing (YouTube, MetaCafe), livecasting (Ustream or Justin.TV), virtual worlds (Kaneva), social gaming (World of Warcraft), social search (Google, Bing, or Ask.com), and instant messaging (Google Talk, Skype, or Yahoo! messenger).

The first social media site was introduced by Geocities in 1994, which allowed users to create their own homepages. The first social networking site, SixDegree.com, was introduced in 1997. Since then, many other social media sites have been introduced, each providing service to millions of people. These individuals form a virtual world in which individuals (social atoms), entities (content, sites, etc.) and interactions (between individuals, between entities, between individuals and entities) coexist. Social norms and human behavior govern this virtual world. By understanding these social norms and models of human behavior and combining them with the observations and measurements of this virtual world, one can systematically analyze and mine social media.

Social media mining is the process of representing, analyzing, and extracting meaningful patterns from data in social media, resulting from social interactions. It is an interdisciplinary field encompassing techniques from computer science, data mining, machine learning, social network analysis, network science, sociology, ethnography, statistics, optimization, and mathematics. Social media mining faces grand challenges such as the big data paradox, obtaining sufficient samples, the noise removal fallacy, and evaluation dilemma.

Social media mining represents the virtual world of social media in a computable way, measures it, and designs models that can help us understand its interactions. In addition, social media mining provides necessary tools to mine this world for interesting patterns, analyze information diffusion, study influence and homophily, provide effective recommendations, and analyze novel social behavior in social media.

1.5 Bibliographic Notes

For historical notes on social media sites and challenges in social media, refer to [81, 173, 141, 150, 115]. Kaplan and Haenlein [141] provide a categorization of social media sites into collaborative projects, blogs, content communities, social networking sites, virtual game worlds, and virtual social worlds. Our definition of social media is a rather abstract one whose elements are social atoms (individuals), entities, and interactions. A more detailed abstraction can be found in the work of [149]. They consider the seven building blocks of social media to be identity, conversation, sharing, presence, relationships, reputation, and groups. They argue that the amount of attention that sites give to these building blocks makes them different in nature. For instance, YouTube provides more functionality in terms of groups than LinkedIn.

Social media mining brings together techniques from many disciplines. General references that can accompany this book and help readers better understand the material in this book can be found in data mining and web mining [120, 280, 92, 174, 51], machine learning [40], and pattern recognition [75] texts, as well as network science and social network analysis [78, 253, 212, 140, 28] textbooks. For relevant references on optimization refer to [44, 219, 228, 207] and for algorithms to [61, 151]. For general references on social research methods consult [36, 47]. Note that these are generic references and more specific references are provided at the end of each chapter. This book discusses non-multimedia data in social media. For multimedia data analysis refer to [49].

Recent developments in social media mining can be found in journal articles in IEEE Transactions on Knowledge and Data Engineering (TKDE), ACM Transactions on Knowledge Discovery from Data (TKDD), ACM Transactions on Intelligent Systems and Technology (TIST), Social Network Analysis and Mining (SNAM), Knowledge and Information Systems (KAIS), ACM Transactions on the Web (TWEB), Data Mining and Knowledge Discovery (DMKD), World Wide Web Journal, Social Networks, Internet Mathematics, IEEE Intelligent Systems, and SIGKDD Exploration. Conference papers can be found in proceedings of Knowledge Discovery and Data Mining (KDD), World Wide Web (WWW), Association for Computational Linguistics (ACL), Conference on Information and Knowledge Management (CIKM), International Conference on Data Mining (ICDM), Internet Measuring Conference (IMC), International Con-

ference on Weblogs and Social Media (ICWSM), International Conference on Web Engineering (ICWE), Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD), The European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Data Basis (ECML/PKDD), Web Search and Data Mining (WSDM), International Joint Conferences on Artificial Intelligence (IJCAI), Association for the Advancement of Artificial Intelligence (AAAI), Recommender Systems (RecSys), Computer-Human Interaction (CHI), SIAM International Conference on Data Mining (SDM), Hypertext (HT), and Social Computing Behavioral-Cultural Modeling and Prediction (SBP) conferences.

Table 1.1: List of Websites

Amazon	Flickr	Facebook	Twitter
BlogCatalog	MySpace	Last.fm	Pandora
LinkedIn	Reddit	Vimeo	Del.icio.us
StumbleUpon	Yelp	YouTube	Meetup

1.6 Exercises

1. Discuss some methodologies that can address the grand challenges of social media.
2. What are the key characteristics of social media that differentiate it from other media? Please list at least two with a brief explanation.
3. What are the different types of social media? Name two, and provide a definition and an example for each type.
4. (a) Visit the websites in Table 1.1 (or find similar ones) and identify the types of activities that individuals can perform on each one.
(b) Similar to questions posed in Section 1.3, design two questions that you find interesting to ask with respect to each site.
5. What marketing opportunities do you think exist in social media? Can you outline an example of such an opportunity in Twitter?
6. How does behavior of individuals change across sites? What behaviors remain consistent and what behaviors likely change? What are possible reasons behind these differences?
7. How does social media influence real-world behaviors of individuals? Identify a behavior that is due to the usage of, say, Twitter.
8. Outline how social media can help NGOs fulfill their missions better in performing tasks such as humanitarian assistance and disaster relief.
9. Identify at least three major side effects of information sharing on social media.

10. Rumors spread rapidly on social media. Can you think of some method to block the spread of rumors on social media?